

CS 314 Principles of Programming Languages

Lecture 24: CUDA Programming

Prof. Zheng Zhang



Rutgers University


December 5, 2018

Class Information

- Project 3 released: deadline 12/12.
- Homework 8 released: deadline 12/10.
- Final exam coverage:
Lecture 1 - 21, HW 1-8, Recitation 1-12, Corresponding Book Chapters.
- Final exam: 12/19 4pm — 7pm.

GPU Programming

- General Purpose Graph Processing Unit (GPU)
- Reflects the trend of multi-core scaling in post Moore's law era
- Application domain: machine learning, scientific simulation, weather prediction, computer vision, bioinformatics, and etc.



Difference between CPU and GPU Programming

CPU Program

Add vector A and vector B to vector C.

```
void add_vector (float *A,  
                float *B,  
                float *C,  
                int N)  
{  
    for ( int i = 0; i < N; i++ )  
        C[i] = A[i] + B[i];  
}
```

```
void main ( ) { ...  
    add_vector (A, B, C, N);  
...}
```

GPU Program

Add vector A and vector B to vector C.

```
_global_ void add_vector (float *A,  
                          float *B,  
                          float *C,  
                          int N)  
{  
    if ( tid < N )  
        C[tid] = A[tid] + B[tid];  
}
```

```
void main ( ) { ...  
    add_vector <<<dimGrid, dimBlock>>> (A, B, C, N);  
...}
```

Programming in CUDA

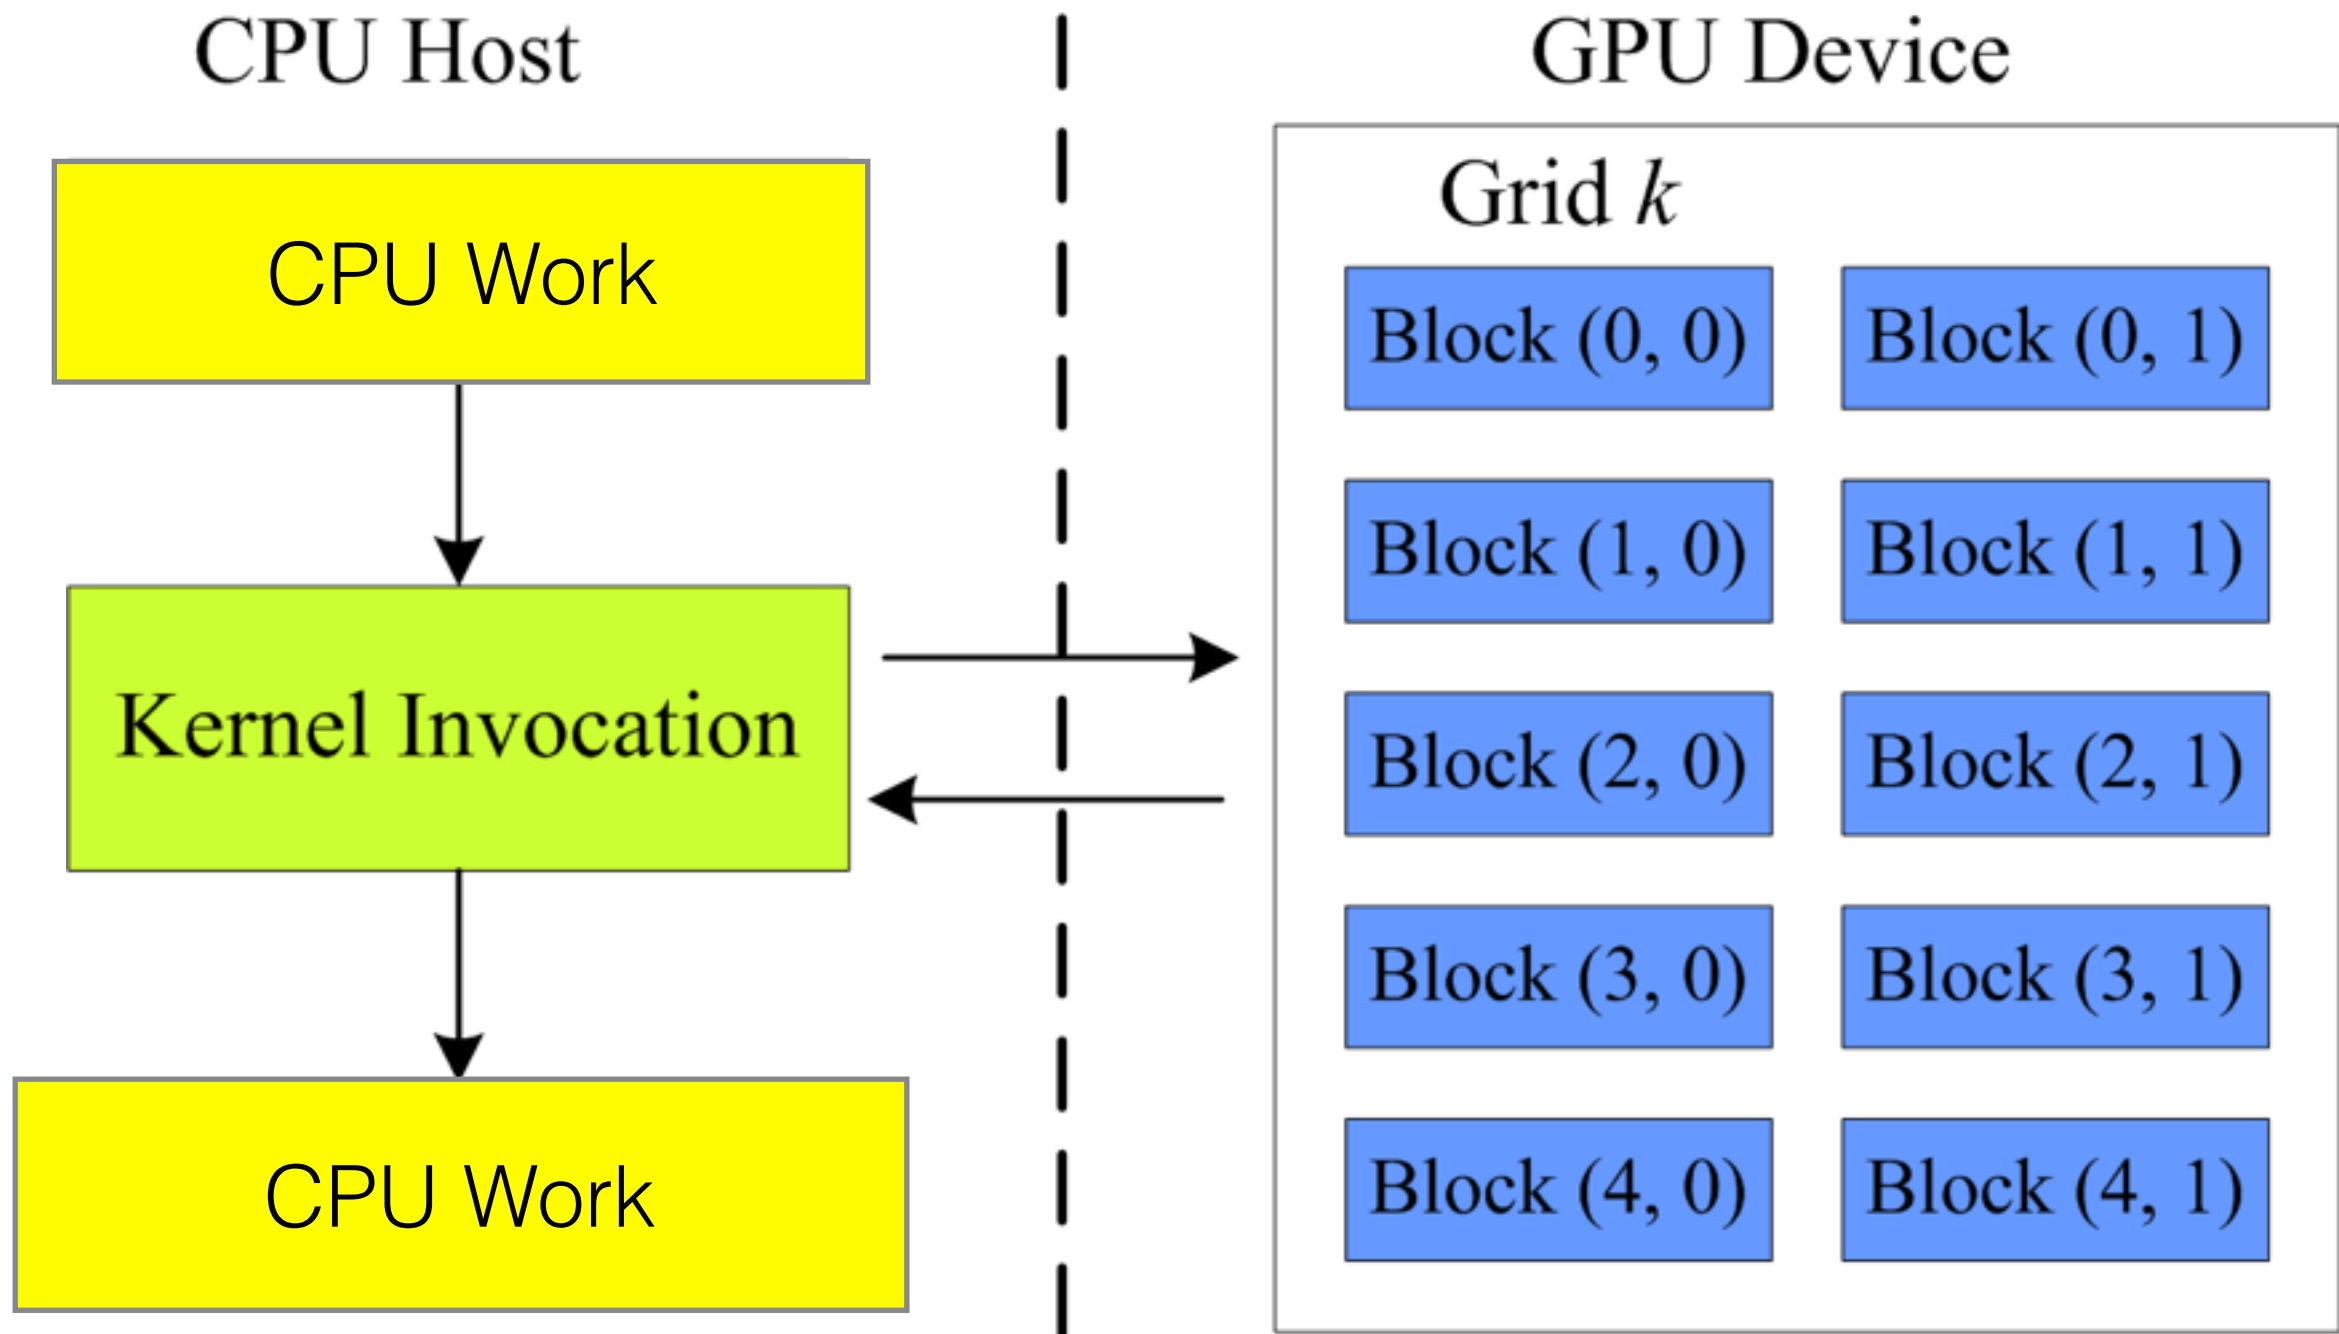
Device Code (GPU):

```
_global_ void add_vector (float *A, float *B, float *C, int N){  
    tid = blockDim.x * blockIdx.x + threadIdx.x;  
    if (tid < N)  
        C[tid] = A[tid] + B[tid];  
    }
```

Host Code (CPU):

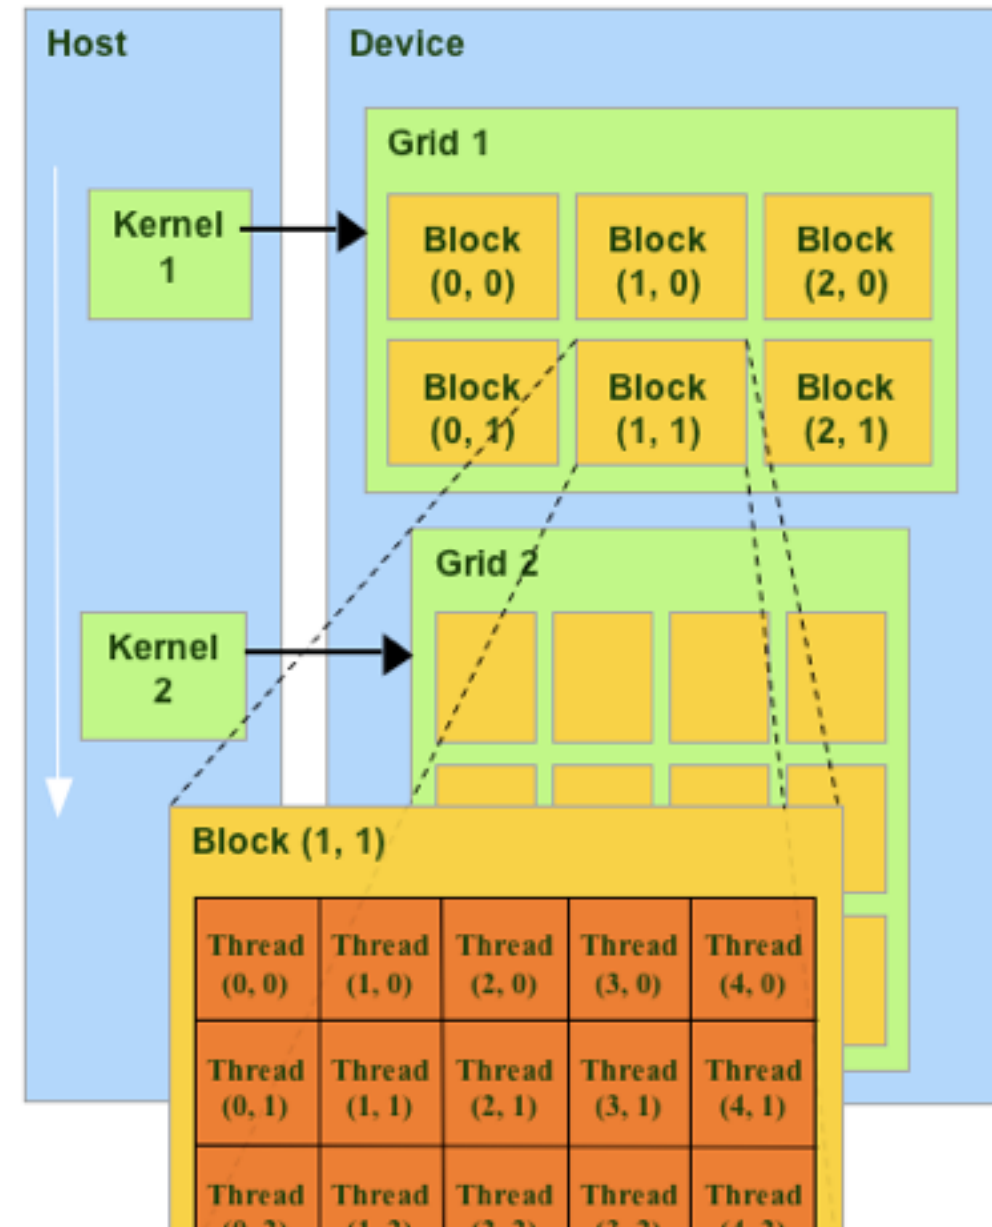
```
void main ( ) {  
    ... //allocate memory on GPU and initialization  
    add_vector <<<dimGrid, dimBlock>>> (A, B, C, N);  
    ... // free memory on GPU  
    ...}
```

Programming in CUDA



Thread View

- A kernel is executed as a grid of thread blocks
- A thread block is a batch of threads that can cooperate with each other by:
 - Synchronizing their execution
 - Sharing data through a low latency scratch-pad memory (named “shared memory”)
- Two threads from two different blocks cannot communicate through the scratch-pad memory



Single Instruction Multiple Data (SIMD)

- **Flynn's taxonomy for modern processors**
 - Single Instruction Single Data (SISD)
 - Single Instruction Multiple Data (SIMD)
 - Multiple Instruction Single Data (MISD)
 - Multiple Instruction Multiple Data (MIMD)

$$C[tid] = A[tid] + B[tid];$$

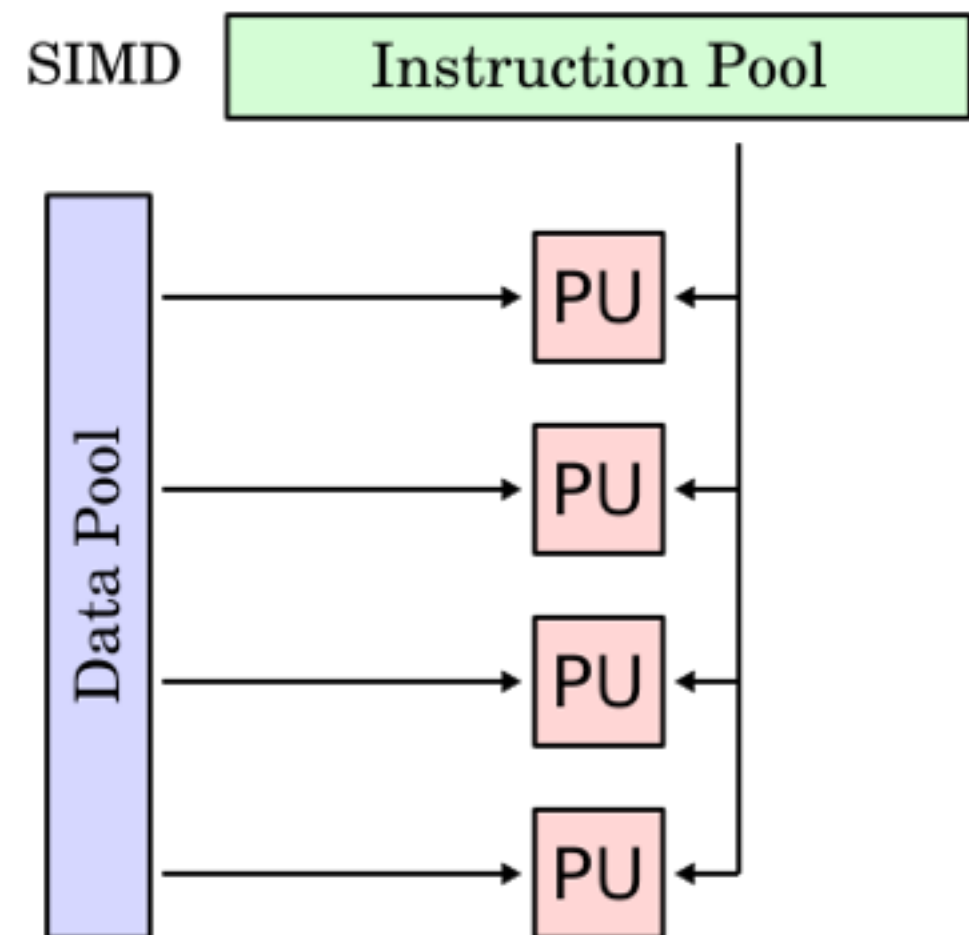
thread 0: $C[0] = A[0] + B[0];$

thread 1: $C[1] = A[1] + B[1];$

....

thread 15: $C[15] = A[15] + B[15];$

Each PU is a processing unit



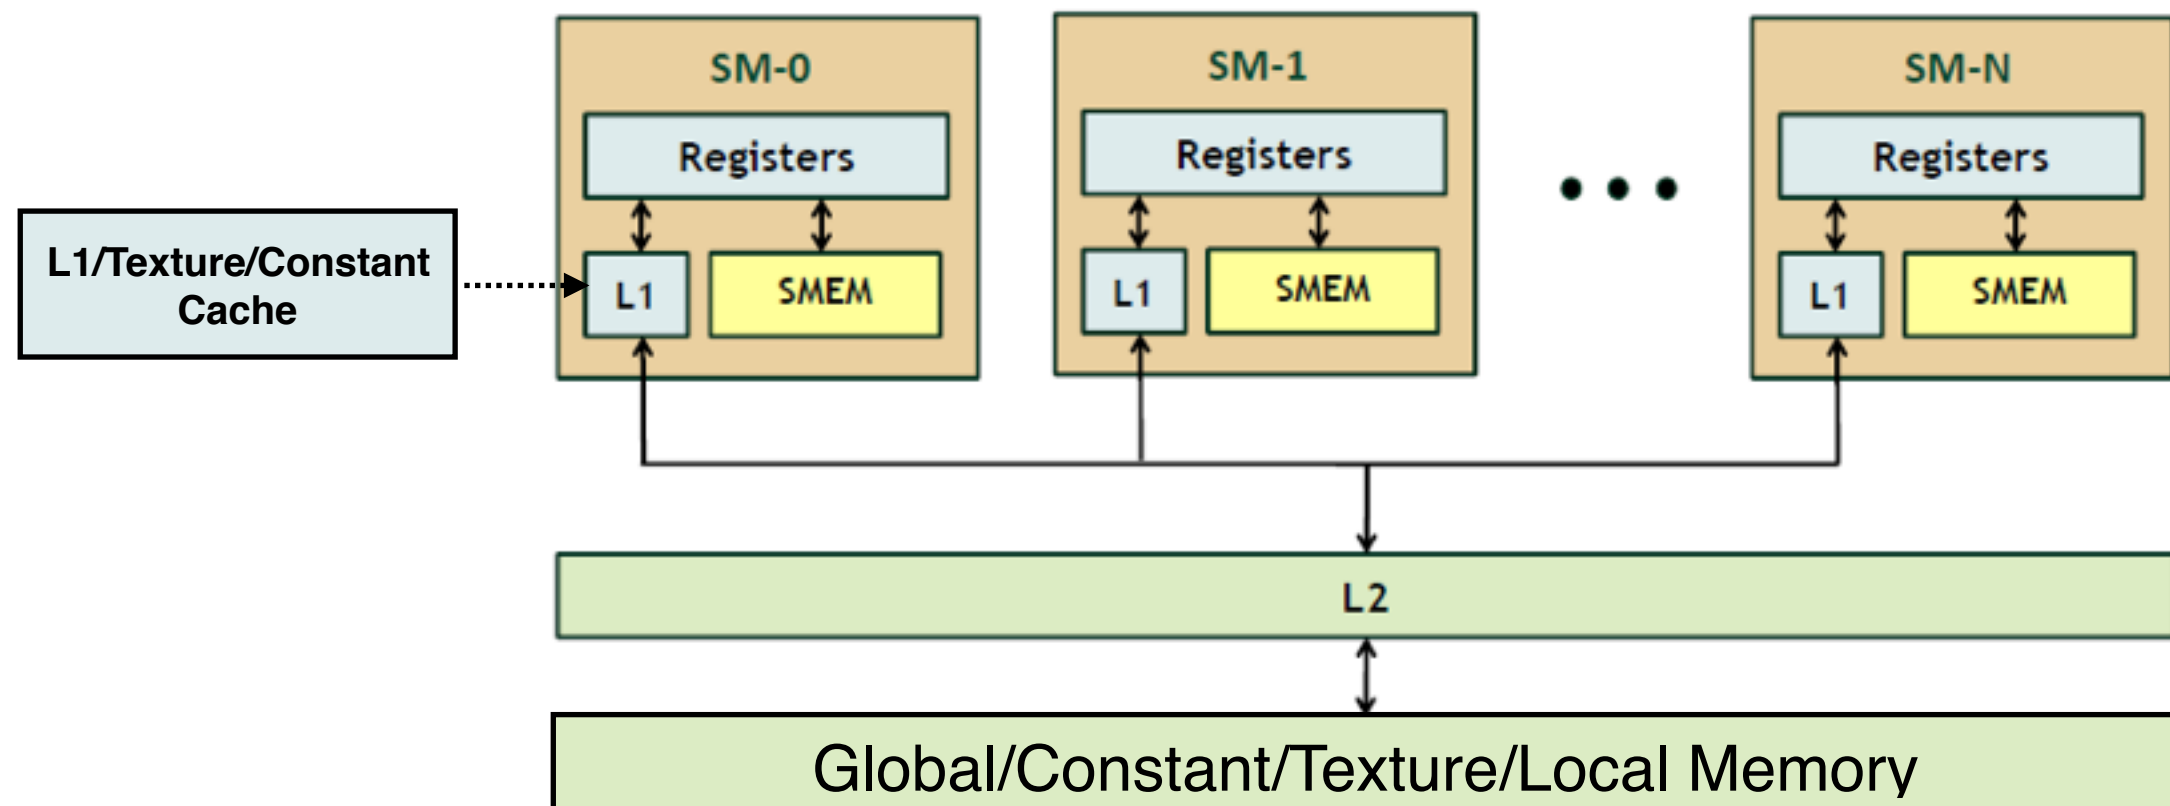
Single Instruction Multiple Thread (SIMT)

- CUDA adopts Single Instruction Multiple Thread (SIMT) model
 - The machine model is organized into multiple SIMD units (warps)
 - Within each SIMD unit, the threads execute the same instruction
 - Across SIMD units, the threads execute different instructions
 - The software model hides the hardware complexity
 - SIMT allows specification of independent branching behavior, while SIMD does not
 - SIMD exposes the width to software writers, while SIMT does not

SIMT enables programmers to write thread-level parallel code for independent and scalar threads, even though the underlying hardware requires co-ordination between threads.

CUDA Memory Model


- On Chip Memory
 - Registers and Shared Memory (SMEM)
 - Instruction Cache, Constant Cache, Texture/L1 Cache
 - Small latency, i.e., ~ 20 cycles or less
- Off Chip Memory
 - L2 (last level cache), ~ 400 -1000 cycles latency
 - Global/Constant/Texture/Local Memory



CUDA Synchronization

- **Barrier Synchronization**

- Within a thread block: `__syncthreads()`
- All computation before the barrier must complete before any computation after the barrier begins
- Barriers divide computation into phases



CUDA Synchronization

- **Compare-and-Swap (CAS)**
 - atomicCAS is provided
 - atomicAdd, atomicSub, atomicInc, and etc
 - built-in atomic functions can be used to implement other synchronization primitives such as locks, mutex, and monitors.


Simple Implementation of Matrix Transpose

- Each thread is in charge of one cell in the matrix

```
__global__ void transpose_naive(float * outputData,  
                                float * inputData,  
                                int width,  
                                int height)  
{  
    unsigned int xIndex = blockDim.x * blockIdx.x + threadIdx.x;  
    unsigned int yIndex = blockDim.y * blockIdx.y + threadIdx.y;  
  
    if (xIndex < width && yIndex < height) {  
        unsigned int index_in = xIndex + width * yIndex;  
        unsigned int index_out = yIndex + height * xIndex;  
        outputData[ index_out ] = inputData[ index_in ];  
    }  
}
```

Case Study I


- **Matrix Transpose**



Simple Implementation of Matrix Transpose

- Each thread is in charge of one cell in the matrix


```
__global__ void transpose_naive(float * outputData,  
                                float * inputData,  
                                int width,  
                                int height)  
{  
    unsigned int xIndex = blockDim.x * blockIdx.x + threadIdx.x;  
    unsigned int yIndex = blockDim.y * blockIdx.y + threadIdx.y;  
  
    if (xIndex < width && yIndex < height) {  
        unsigned int index_in = xIndex + width * yIndex;  
        unsigned int index_out = yIndex + height * xIndex;  
        outputData[ index_out ] = inputData[ index_in ];  
    }  
}
```



Global Memory Coalescing

- **Data is fetched as a contiguous memory chunk**
 - A cache line or a cache block, typically 128 byte for GPUs
 - The purpose is to utilize wide DRAM burst for maximum memory level parallelism (MLP)s
- **A thread warp is ready when all its threads' operands are ready**


$P[] = \{0, 1, 2, 3, 4, 5, 6, 7\}$



a cache block.

Coalesced Accesses

$P[] = \{0, 5, 1, 7, 4, 3, 6, 2\}$



a cache block.

Non-Coalesced Accesses

Simple Implementation of Matrix Transpose


- Each thread is in charge of one cell in the matrix

```
__global__ void transpose_naive(float * outputData,  
                                float * inputData,  
                                int width,  
                                int height)  
{  
    unsigned int xIndex = blockDim.x * blockIdx.x + threadIdx.x;  
    unsigned int yIndex = blockDim.y * blockIdx.y + threadIdx.y;  
  
    if (xIndex < width && yIndex < height) {  
        unsigned int index_in = xIndex + width * yIndex;  
        unsigned int index_out = yIndex + height * xIndex;  
        outputData[ index_out ] = inputData[ index_in ];  
    }  
}
```

coalesced or not?

Matrix Transpose

- Let's try something else!



Matrix Transpose

- Improved Implementation

```
__global__ void transpose(float *outputData, float *inputData, int width, int height){  
    __shared__ float block[BLOCK_DIM][BLOCK_DIM + 1];
```

Input: 1 Million Matrix Elements

Naive Implementation

Performance: Throughput = 8.0226 GB/s, Time = 0.97381 ms

Improved implementation


Performance: Throughput = 11.5080 GB/s, Time = 0.67887 ms

Code: /ilab/users/zz124/cs515/samples/6_Advanced/transpose

```
    (yIndex < height) && (xIndex < width)) {  
        unsigned int index_out = yIndex * height + xIndex;  
        outputData[index_out] = block[threadIdx.x][threadIdx.y];  
    }  
}
```

Case Study II

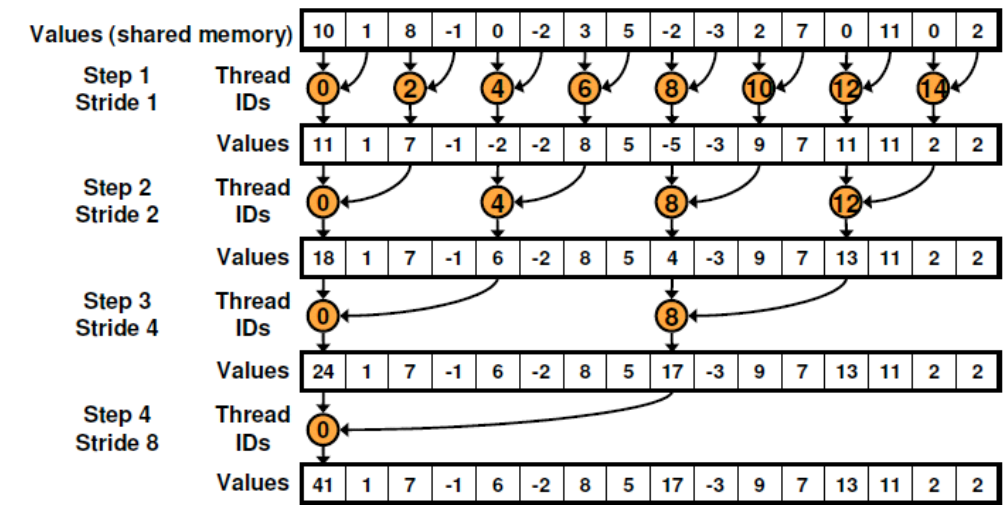
- Parallel Reduction



Case Study II


• Parallel Reduction Code

```
tid = threadIdx.x;  
  
for ( s=1; s < blockDim.x; s *= 2 ) {  
    if ( tid % (2*s) == 0 ) {  
        sdata[ tid ] += sdata[ tid + s ];  
    }  
    __syncthreads();  
}
```



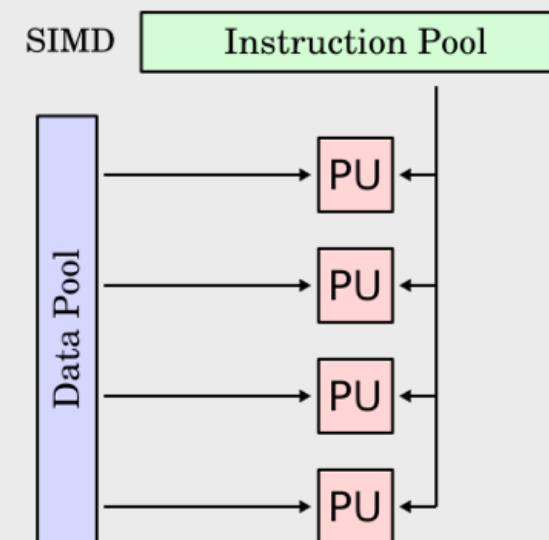
Control Divergence

- Different runtime execution path lead to processor underutilization




```
if (A[tid]) {  
    do some work;  
} else {  
    do nothing;  
}
```

Recall that in SIMD model, only one instruction can be fetched and executed at one time. If any thread does not execute that instruction, the corresponding PU will be idle.



Parallel Reduction

- Optimize Control Divergence



First Version



Second Version

Parallel Reduction

- **Optimize Control Divergence**

Reduction on 16 Million Elements

Implementation 1:

Performance: Throughput = 1.5325 GB/s, Time = 0.04379 s

Implementation 2:

Performance: Throughput = 1.9286 GB/s, Time = 0.03480 s


First Version




Second Version

Parallel Reduction

- Improve Coalescing



Second Version

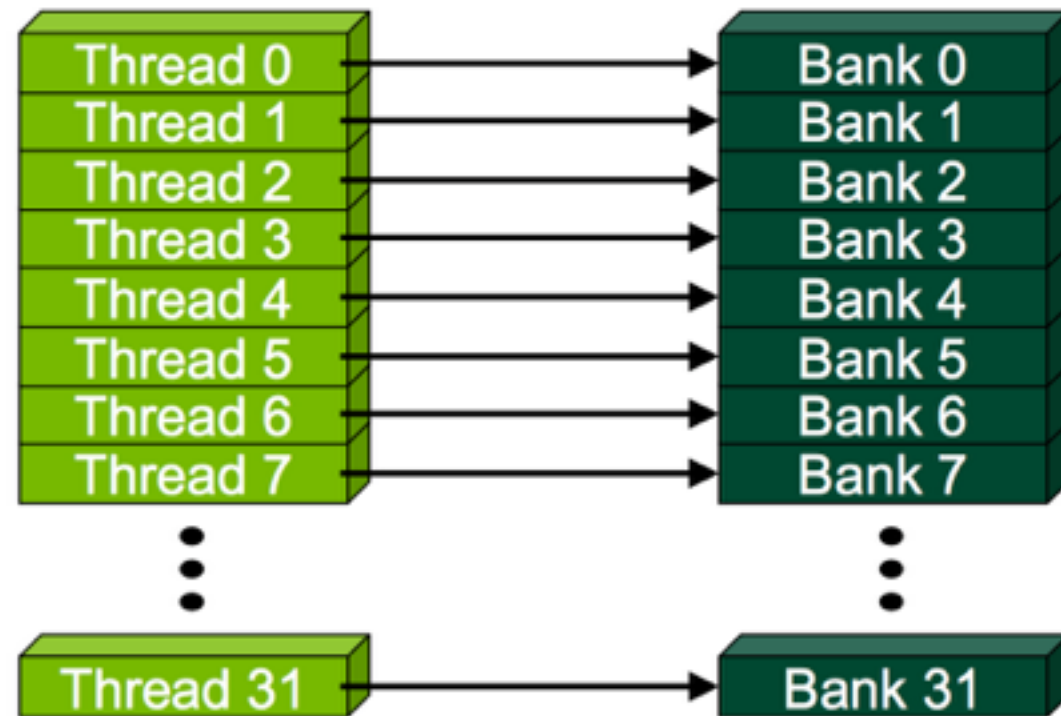


Third Version

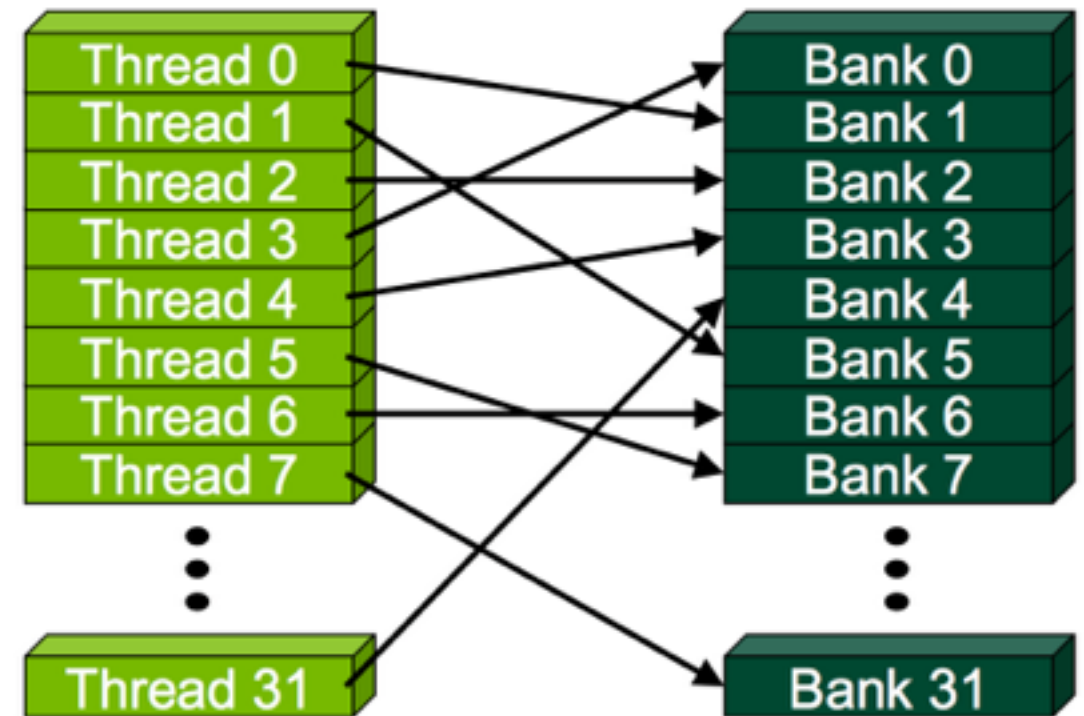
Bank Conflicts

- **Shared Memory in CUDA Typically 16 or 32 banks**
 - Successive 32-bit words are assigned to successive banks
 - A fixed stride access may cause bank conflicts
 - Maximizing bank level parallelism is important

- No Bank Conflicts



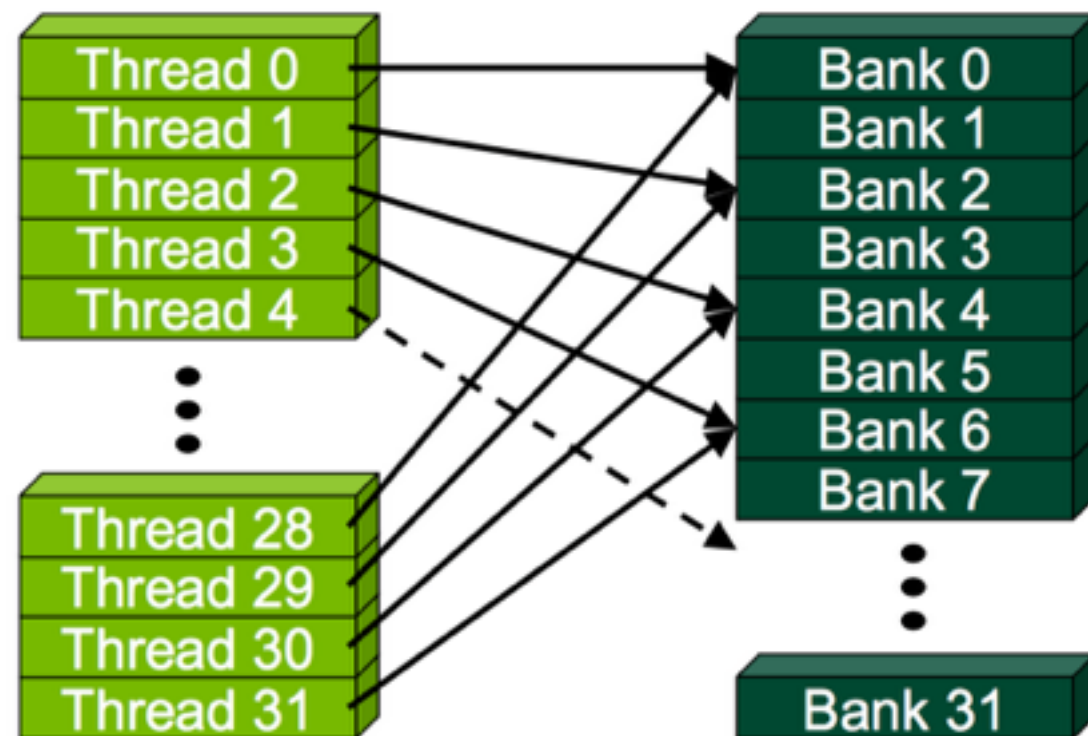
- No Bank Conflicts



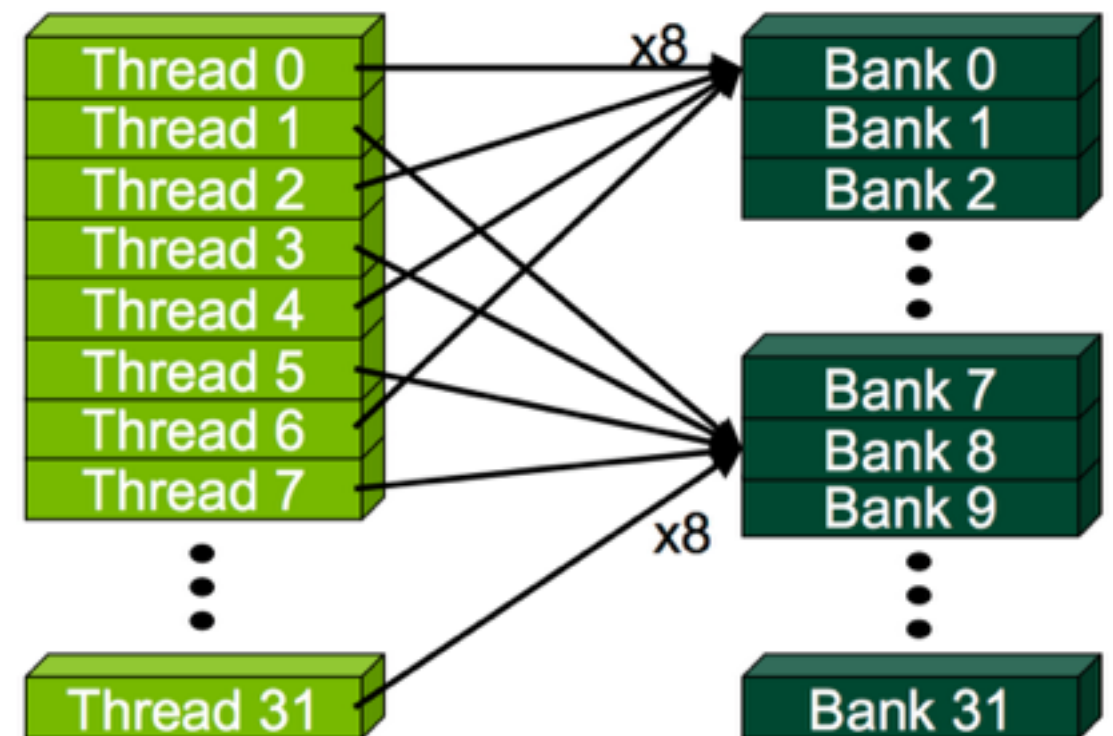
Bank Conflicts

- **Shared Memory in CUDA Typically 16 or 32 banks**
 - Successive 32-bit words are assigned to successive banks
 - A fixed stride access may cause bank conflicts
 - Maximizing bank level parallelism is important

- **2-way Bank Conflicts**




- **8-way Bank Conflicts**




Parallel Reduction

- Improve Coalescing and Eliminates Bank Conflicts



Second Version



Third Version

GPU Resources at ilab

- **1 multi-GPU Machine: “atlas.cs.rutgers.edu”**
 - 1x Tesla K40:
12 GB memory, 2880 CUDA cores, CUDA capability 3.5
 - 2x Quadro K4000:
3 GB memory, 768 CUDA cores, CUDA capability 3.0
- **20 single-GPU Machines (listed on the right)**
 - Each has one GT 630
 - 2 GB memory, 192 CUDA cores, CUDA capability 3.0

Check <http://report.cs.rutgers.edu/mrtg/systems/ilab.html> for availability of different GPU machines

Reading

- “Optimization Parallel Reduction in CUDA” by Mark Harris, NVIDIA
http://developer.download.nvidia.com/compute/cuda/1.1-Beta/x86_website/projects/reduction/doc/reduction.pdf
- “Parallel Prefix Sum (Scan) with CUDA”, GPU Gems3, Chapter 39, by Mark Harris, Shubhabrata Sengupta, and John Owens.
http://http.developer.nvidia.com/GPUGems3/gpugems3_ch39.html
- “Performance Optimization” by Paulius Micikevicius, NVIDIA
<http://www.nvidia.com/docs/IO/116711/sc11-perf-optimization.pdf>
- CUDA Programming Guide
<https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html>